

Local Search and Genetic Algorithms

Dr. Demetrios Glinos
University of Central Florida


CAP4630 - Artificial Intelligence

Today

- Local Search
- Hill Climbing
- Simulated Annealing
- Beam Search
- Genetic Algorithms

Classic Search Problems Revisited

- Recall how we first defined a search problem
 - state space
 - successor function
 - start state
 - goal test
 - solution is a *sequence of actions (the plan)* that transforms the start state into a goal state
- We could use tree search algorithms to find optimal solutions
 - e.g., BFS, UCS
- We could also use heuristic algorithms
 - e.g., Greedy, A*




Optimization Problems

- Consider
 - state space too large to exhaustively enumerate
 - no particular start state
 - we don't know if we are in a goal state
 - path to a goal state doesn't matter, as long as we get there
 - optimal solution desirable, but not required
- This is the domain of many optimization problems
 - e.g., manufacturing process optimization, TSP
- Local search methods can often find reasonable solutions

Local Search

- **Recall:** **Tree search** keeps still-viable unexplored alternatives on a fringe
 - ensures completeness, goal is to find any goal state
- **Local search:**
 - An optimization technique for searching a space that cannot be exhaustively enumerated
 - Goal is to find the global maximum
 - Basic idea: **Improve a single option until you can't make it better**
- **Since there is no fringe:**
 - Generally, much faster and uses less memory
 - BUT, incomplete and suboptimal
- **Requirement:** **an evaluation function**
 - to determine what is an improvement





Today

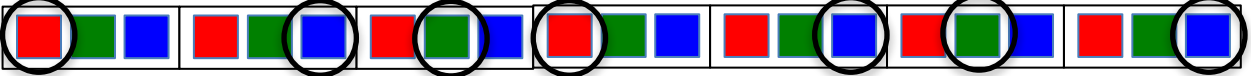
- Local Search
- Hill Climbing
- Simulated Annealing
- Beam Search
- Genetic Algorithms

Hill Climbing


- **Hill climbing (aka greedy local search)**
- **Basic idea:**
 - Use an evaluation (heuristic) function for the states
 - Start anywhere
 - Move in direction of steepest gradient
 - If no neighbor better than current, then stop
- Can get stuck at local maximum
- Completeness and optimality not assured
- Nevertheless, can be useful in practice
 - when a “good” solution is good enough
 - when the search space is relatively well behaved



Example: CSP Min Conflicts

- Map-coloring CSP: 
 - Basic idea
 - Start with a complete assignment with unsatisfied constraints
 - Use operators to reassign variable values
 - Iterate until a solution found or exhaust all possibilities
 - Note: No fringe – work with just one assignment !
 - Iterative Min Conflicts algorithm
 - Variable selection: random choice from among conflicting variables
 - “Min conflicts” value selection: choose value that results in fewest constraint violations
- This is hill climbing with the function: $\text{heuristic } h(n) = \text{total number of constraints violated}$

Hill-Climbing Example: 4-Queens




- **States:**
 - 4 queens, 1 in each column ($4^4 = 256$ total states)
- **Operator:**
 - Move a queen vertically in its column
- **Goal test:**
 - No queen threatens another
- **Heuristic:**
 - $h(n)$ = number of bidirectional attacks

Example: Local Maximum

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

(a)




(b)

source: Fig. 4.3

- Consider: 8-queens problem, where each queen can move only within its column
- case (a): $h = 17$; best moves for each queen are marked
- case (b): a local maximum with $h = 1$ (not a goal state); every possible move has a higher cost

Improving Hill Climbing

- Basic hill climbing is **steepest ascent**
- Stochastic hill climbing
 - choose randomly from among several uphill choices
 - can weight choices by the steepness of the gradient



- First choice hill climbing
 - generate successors until find one that improves cost, and then take it
 - suitable where a state can have many (e.g., thousands) of successors
- Random restart hill climbing
 - perform as many climbs as needed until find an acceptable solution
 - each climb starts at a different random state
 - number of trials needed = $1/p$, where p = prob of success for 1 trial
 - e.g., for 8-queens (column moves only), $p \approx 0.14$, so need roughly 7 trials

Today


- Local Search
- Hill Climbing
- Simulated Annealing
- Beam Search
- Genetic Algorithms

Simulated Annealing

Annealing, in metallurgy, involves heating and slowly cooling a metal, allowing it to form a more regular crystalline structure, which allows the metal to be deformed and to be shaped more easily

Simulated Annealing

- Purpose: to escape local maxima
- How: incorporate a **temperature** parameter into hill climbing
 - at high temperatures: more **exploration**
 - at low temperatures: more **exploitation**
 - use a **cooling schedule** to slowly lower the temperature



Basic Algorithm for Simulated Annealing

Goal: find the global *minimum*



1. Start with a *candidate solution*
2. Calculate its value using some *cost function*
3. Generate a *neighboring solution*
4. Calculate the neighboring solution's cost
5. Compare them:
 - a. If $\text{cost}_{\text{new}} < \text{cost}_{\text{old}}$, accept the new solution
 - b. If $\text{cost}_{\text{new}} > \text{cost}_{\text{old}}$, maybe accept the new solution, depending on the temperature
6. Repeat steps 3 to 5 above, slowly reducing the temperature, until convergence or reach a max limit on iterations

Simulated Annealing Details

- Initial candidate:
 - can be truly random
- Cost function:
 - can be anything appropriate (e.g., total distance traveled for TSP)
- Neighboring solution:
 - must differ from current only slightly (must be a neighbor in the search space)
 - choose randomly among the neighbors
- Always accept a lower cost candidate: $\Delta\text{Cost} = \text{cost}_{\text{new}} - \text{cost}_{\text{old}} < 0$
- Acceptance function for higher-cost candidate:
 - Metropolis-Hastings algorithm: accept if $e^{-\Delta\text{Cost}/\text{Temp}} > \text{Random}(0, 1)$
 - threshold acceptance algorithm: accept if $\Delta\text{Cost} < \text{threshold}$
- Temperature schedule:
 - Start temp at value 1.0 and reduce by α every n iterations
 - Typically, $0.8 < \alpha < 0.99$
 - $\text{Temp}_i = \alpha^i$

Example: Simulated Annealing


Sweet spot: Combinatorial optimization, such as Traveling Salesperson Problem (TSP)



Search space for 20-city TSP is $20! = 1.55 \times 10^{25}$

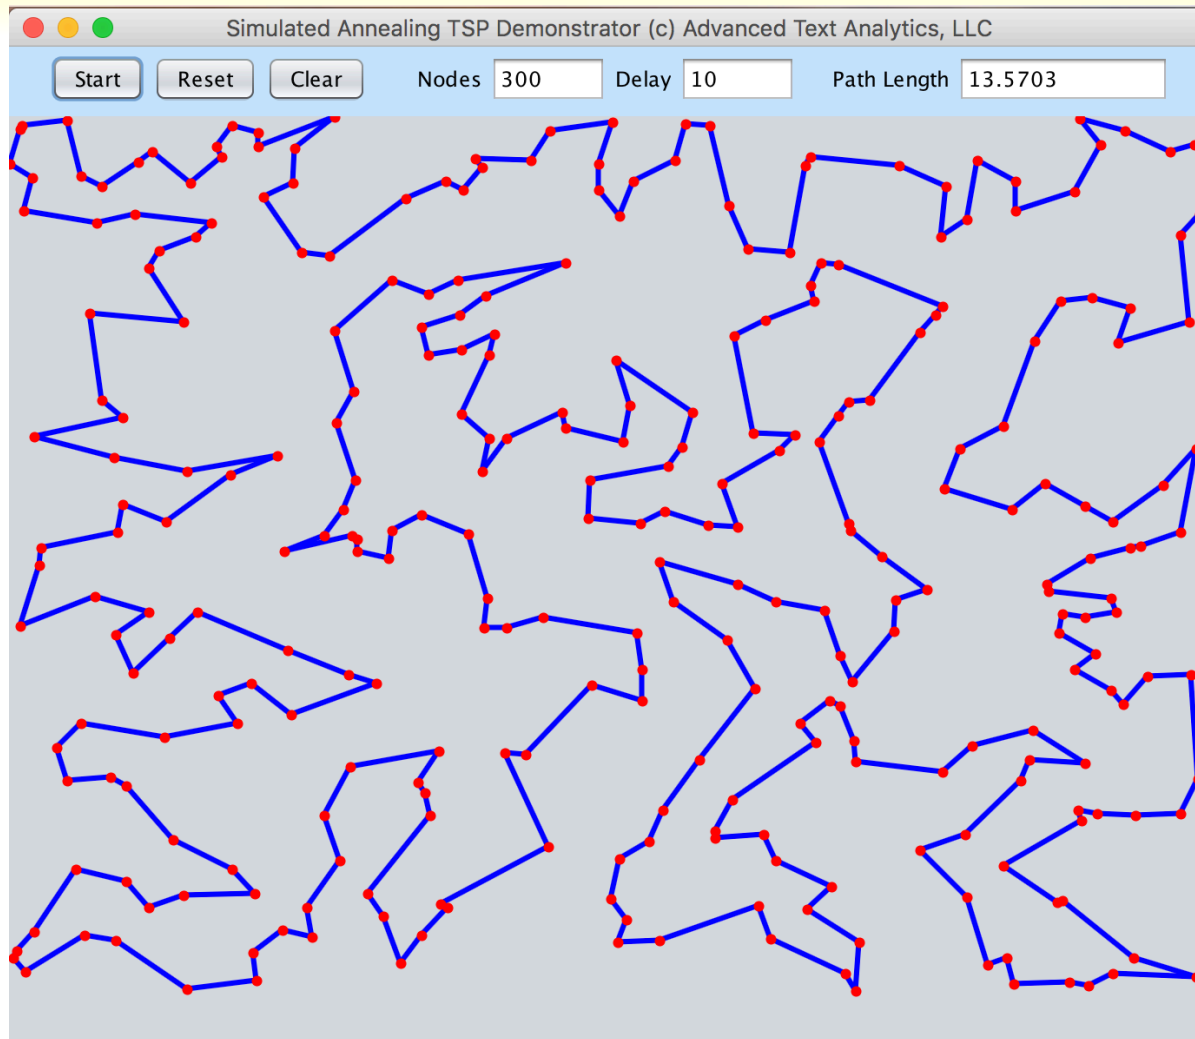
Demo: Simulated Annealing TSP

- Random locations for nodes
- Random initial sequence
 - complete Hamiltonian circuit
- Neighboring solution selection:
 - choose 2 random nodes
 - reverse the path between them and tie into full path
 - compare total cost for revised sequence to current cost
- Cost function: total Euclidean path length
- Cooling rate: $\alpha = 0.9999$



before: solid arrows
after : dashed arrows
cuts as shown

Demo: Simulated Annealing TSP



demo: SimulatedAnnealing.jar


Today

- Local Search
- Hill Climbing
- Simulated Annealing
- **Beam Search**
- Genetic Algorithms


Beam Search

- Like local greedy search, except
 - start with K random initial states
 - at each step, generate all successors of all K states
 - but keep only the K “best” states for next iteration

Local greedy Search




Beam Search



Q: How is this different from having K random restarts?

Improving Beam Search

- **Basic beam search**
 - choose the "best" k successors from the pool of candidate successors
 - can suffer from the lack of diversity among the k states
 - population can be concentrated in a local region of the state space
- **Stochastic beam search**
 - choose a random set of k successors
 - weight each choice according to its value
 - analogy to natural selection
 - population at next step consists of some subset of the successors (offspring) of the current state (population organism)




Today

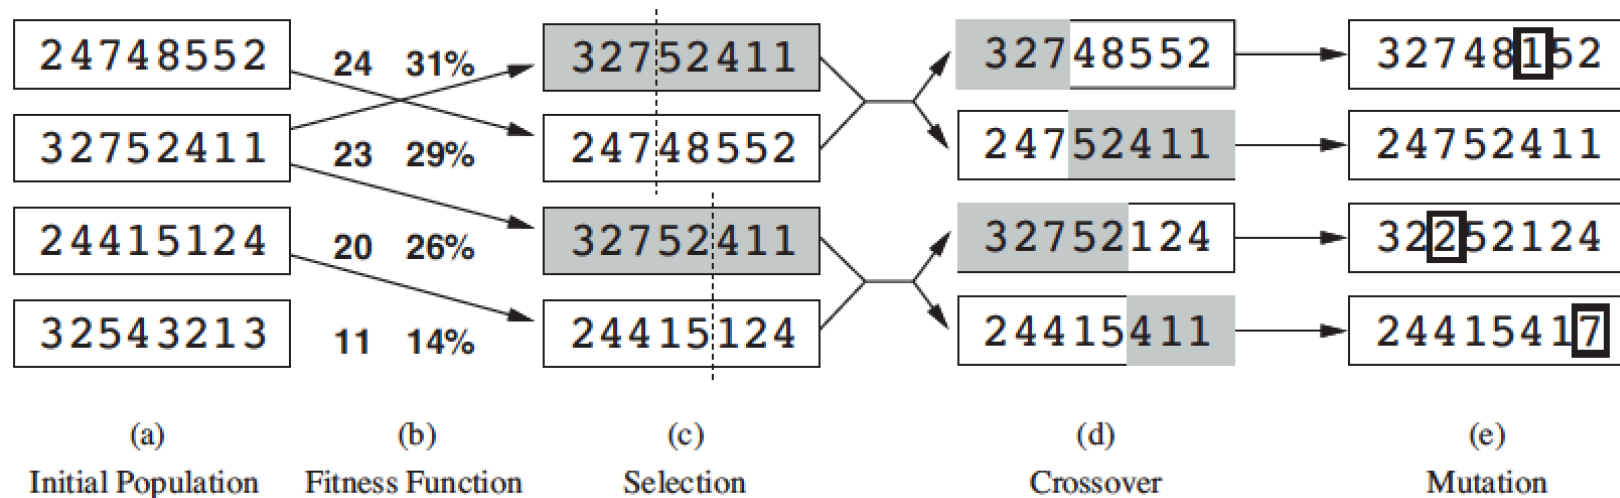
- Local Search
- Hill Climbing
- Simulated Annealing
- Beam Search
- Genetic Algorithms

Genetic Algorithms

- Based on genetics in biology
 - Chromosomes encode traits
 - Offspring inherit “chromosomes” from both parents
 - Mutations can occur
 - Some mutations more successful than others
- Genetic Algorithms
 - encode candidate solutions as a population of *chromosomes*
 - use genetics-inspired operators to generate offspring
 - selection – of parents
 - crossover – inheriting from both parents
 - mutation – random changes in components of candidate solutions
 - evaluate offspring according to a fitness function
 - replace the population with the next generation



GA Details




- **Chromosomes** are generally sequences of symbols, that can represent
 - solutions of complex functions
 - finding sequences of amino acids that will fold to a desired 3-D protein structure
 - sequences of motion for a robot crawler
 - whether or not all the same length depends on the problem

GA Details


- **Selection**
 - **fitness-proportionate** - the more fit a chromosome, the more likely it will be selected as a parent
 - generally done **with replacement** (same chromosome can be parent to > 1)
- **Crossover**
 - usually at 1 randomly-chosen point, with uniform probability p_c
 - point can be different on each chromosome
 - multi-point crossover also possible
- **Mutation**
 - random change to a gene in each chromosome, with probability p_m
 - can restrict to 1 or >1 mutation per chromosome
- **Fitness function**
 - custom for each problem
- **Next Generation**
 - can select top n most fit chromosomes
 - can propagate **elites** (the top few always make it to the next generation)

GA Demo: Robot Crawler




- State space
 - discretized
 - green upper arm: ± 45 degrees from horizontal
 - red lower arm: ± 45 degrees from perpendicular to green arm

GA Demo: Robot Crawler



- Chromosome
 - start with green horizontal and red vertical, touching ground
 - random sequence of NSEW movements
 - do not need all to be same length
 - must validate for
 - movement out of range
 - loops, including stutter


GA Demo: Robot Crawler



- Fitness

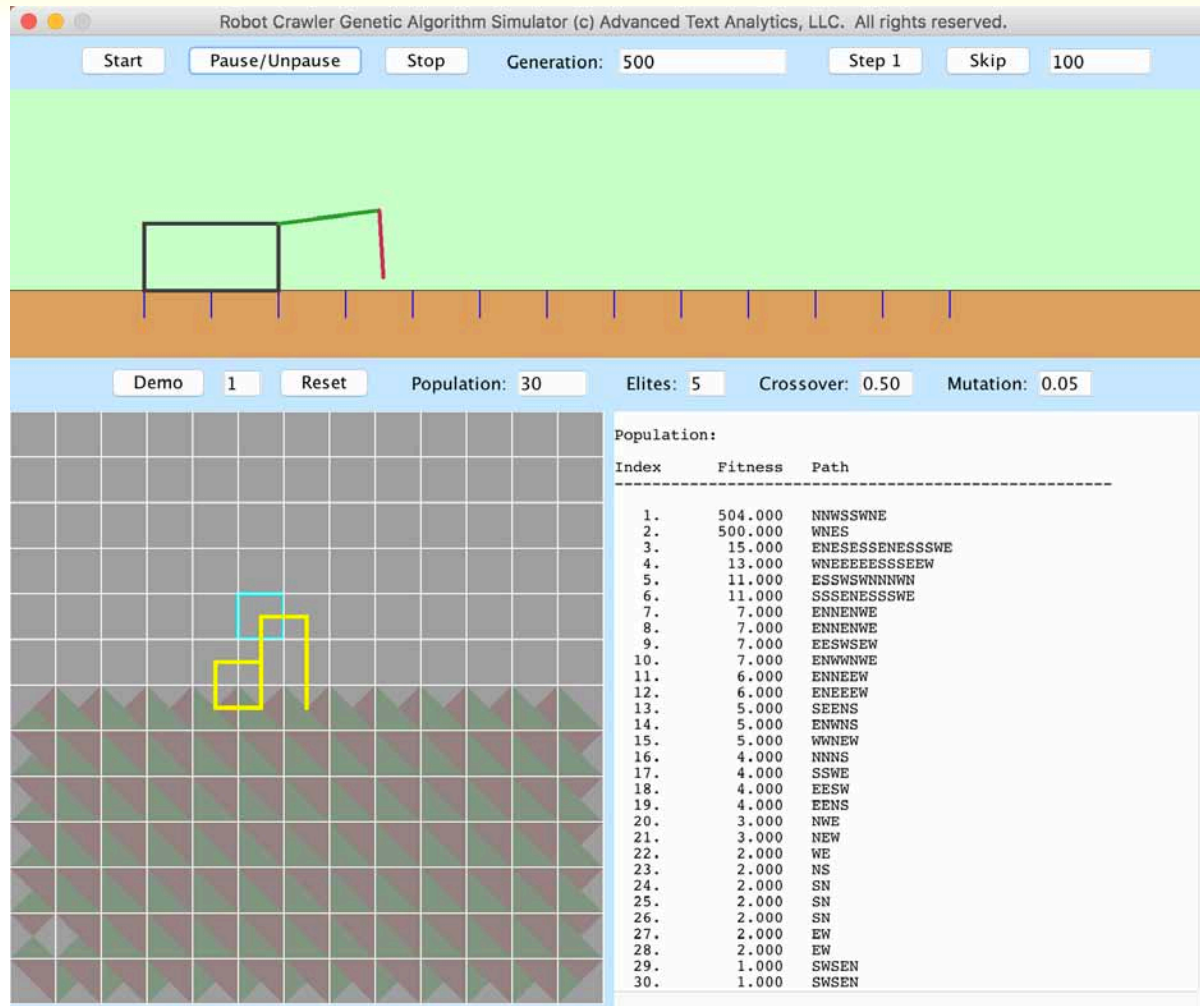
- wish to reward forward motion +500, backward -500, within a loop
- “forward motion” defined to be left corner to the right of previous spot with right corner on ground
- also reward longer sequences, so +1 for every element of sequence outside of a loop

GA Demo: Robot Crawler



- Crossover
 - Single point, must validate
- Mutation
 - Single point, must validate
- Population
 - Start with 30
 - Use elites


GA Demo: Robot Crawler



demo: CrawlGA



GA Problem: Rover Path Planner

- Consider a planetary rover
 - current position known
 - goal position determined externally
 - sensors on board
 - obstacles on path to goal position
- Goal
 - find a safe path to goal position




Path Planner Design

- State space
 - discretized
 - use Cartesian grid
 - could also use hexagonal grid, etc.
- Chromosome
 - sequence of moves
 - must avoid obstacles
- Move
 - vector representation: direction + length
 - 8 directions
 - endpoint must "snap" to center of a grid cell

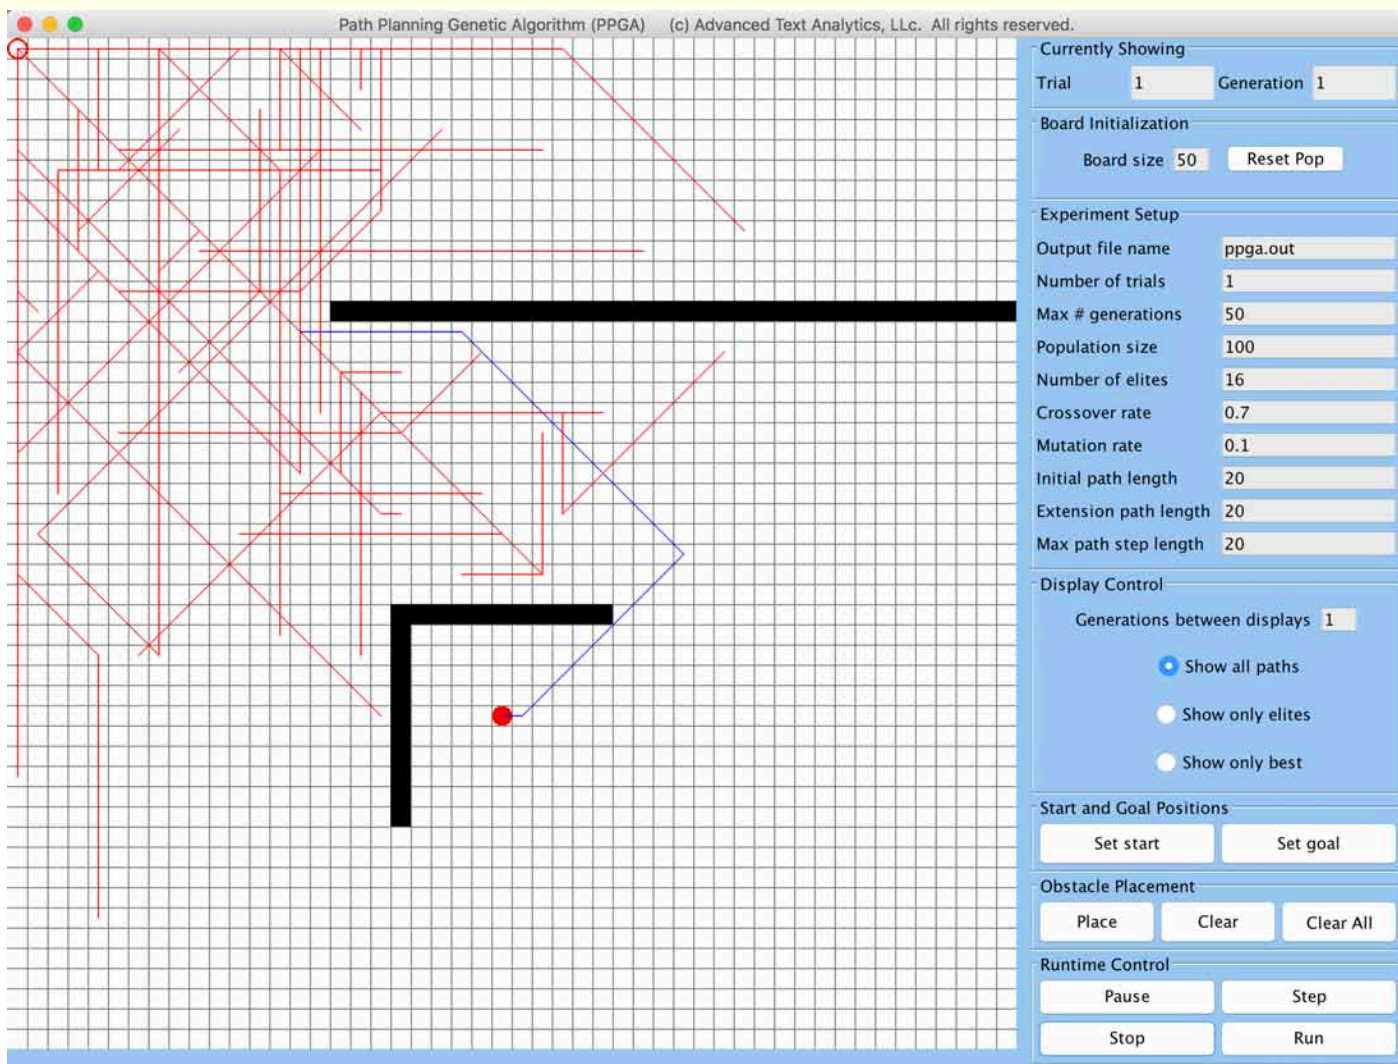


Path Planner Operators

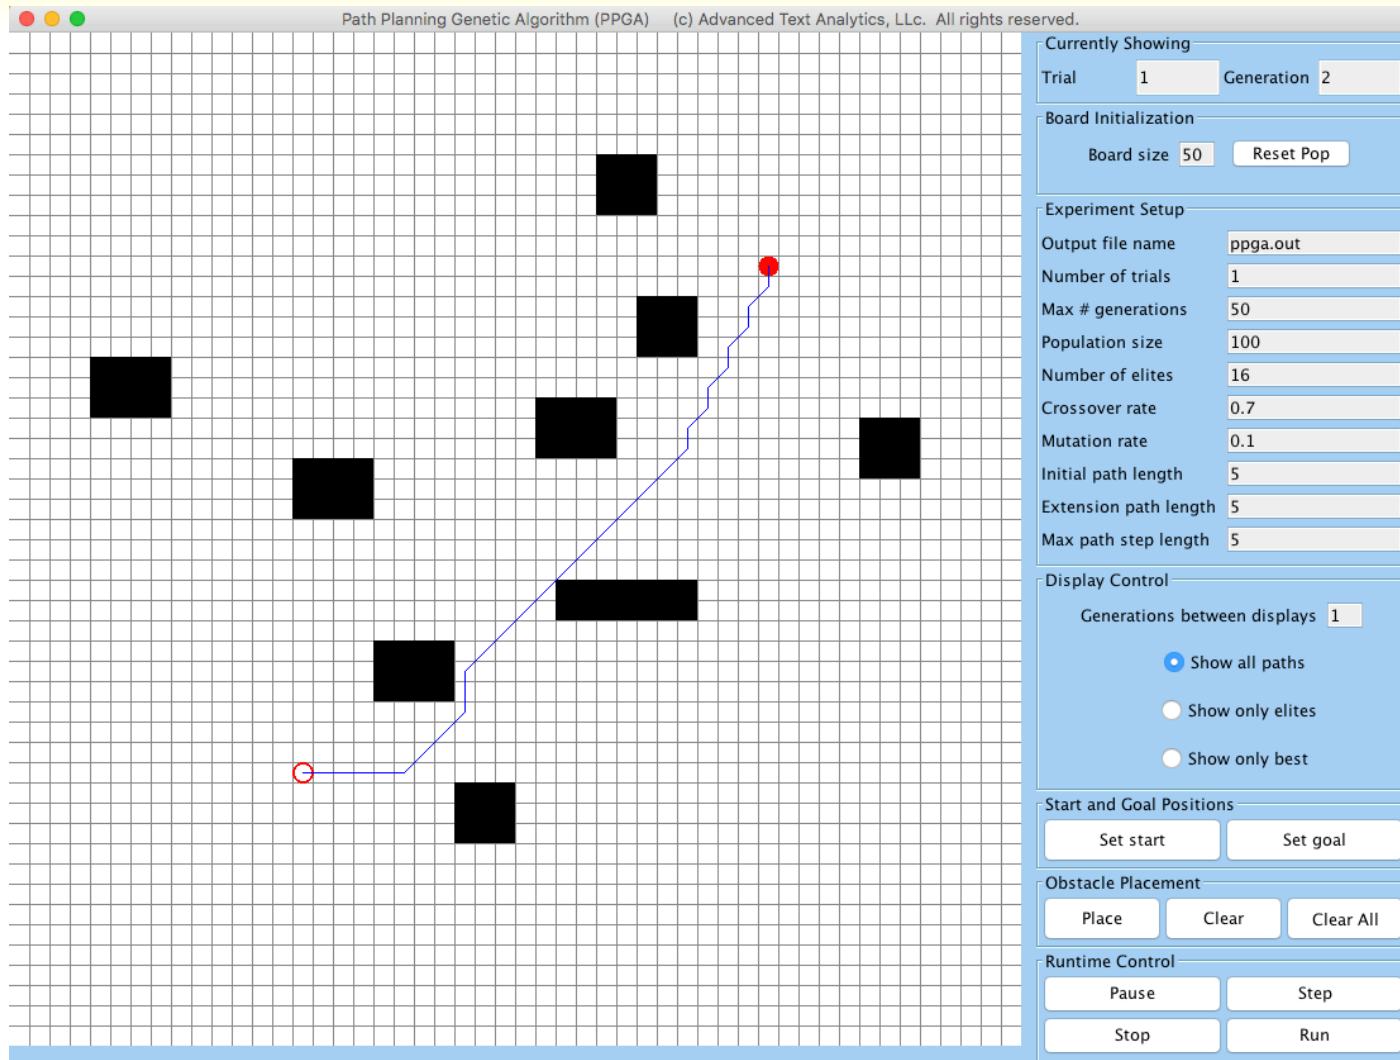
- **Mutation**
 - add random number (bounded) of "legs" to current path
 - each leg can be in 1 of 8 possible directions
 - each leg can be a random length (bounded)
 - prune path if bumps into wall or goes off the board
- **Crossover**
 - single point
 - prune both paths as appropriate
- **Fitness function**
 - use heuristic: Euclidean distance to goal
- **Tuning parameters**
 - crossover rate, random mutation and crossover parameters



GA Demo: Rover Path Planner




GA Demo: Rover Path Planner




demo: PPGA.jar


Improving the Path Found



Cut out loops



Straighten out bights



Cut corners

Note: Can also choose key points in the path and fit a smooth curve (e.g., spline)